

Alast Samimi-Darzi

Data Science Capstone Final Report

The George Washington University

November 2025

AI implemented in Healthcare

Abstract

Racial health disparities in the United States have long been documented, but the COVID-19 pandemic amplified them to unprecedented levels. At the same time, artificial intelligence (AI) and machine-learning-based diagnostic tools became increasingly integrated into healthcare systems. This project examines whether modern predictive models reproduce or mitigate racial disparities by combining epidemiological analysis with algorithmic fairness evaluation.

Using CDC COVID-19 weekly mortality data (2020–2023) and BRFSS (2020–2022) chronic disease survey data, I visualize population-level disparities and train a logistic regression model to predict heart disease. Model performance is evaluated overall and stratified by race using AUC, accuracy, and F1 metrics.

Results show that disparities shift substantially depending on age band: NH/Black, AIAN, and Hispanic adults have the highest midlife (50–64) mortality, while aggregated “all adults” statistics are dominated by the large older NH/White population. The predictive model performs well overall ($AUC \approx 0.82\text{--}0.85$) but demonstrates lower discrimination for NH/Black and AIAN groups. These findings reveal that while AI does not create bias, it reflects structural inequities embedded in health data—underscoring the need for fairness auditing in clinical AI deployment.

Introduction

Through the COVID-19 CDC dataset, we are able to identify exposed and amplified long-standing racial disparities in U.S. health outcomes. Structural factors—such as differential exposure risk, occupational segregation, access to medical care, baseline chronic disease burden, and environmental stressors—produced sharp inequalities in hospitalization and mortality.

At the same time, the healthcare system accelerated the adoption of AI-driven diagnostic and decision-support tools. As these systems influence triage, risk scoring, and treatment recommendations, understanding whether AI reduces or reinforces disparities is essential.

This project examines a central question:

Does AI in healthcare reduce or exacerbate racial disparities in access and outcomes?

To answer this, the project integrates:

1. Epidemiological analysis of COVID-19 mortality using CDC weekly data.
2. Chronic disease and behavioral risk analysis using BRFSS.
3. Predictive modeling for heart disease using logistic regression.
4. Fairness evaluation of accuracy, AUC, and F1 score across racial groups.
5. Two Tableau dashboards synthesizing public health disparities and predictive model performance.

This mixed-methods approach connects population-level disparities to the algorithmic behavior of an AI system designed to predict clinical outcomes.

Literature Review and Data Sources

Public health literature consistently reports severe racial inequities in COVID-19 outcomes. Black, Hispanic, and Indigenous communities experienced higher infection, hospitalization, and mortality rates due to structural determinants such as:

- Essential-worker exposure
- Reduced access to early testing
- Under-insurance and under-treatment
- Higher prevalence of comorbidities related to environmental and socioeconomic stressors

AI fairness research similarly finds that machine-learning models trained on biased or incomplete datasets tend to underperform for marginalized groups—especially when training data reflects structural inequities (Obermeyer et al., 2019; Chen et al., 2021).

CDC COVID-19 Weekly Mortality Data (2020–2023)

- Weekly deaths per 100k, stratified by race and age group
- Aggregated to annual race-specific trends for Tableau visualization
- Used to examine disparities for all adults (18+) and 50–64-year-old adults

BRFSS 2020–2022 (Behavioral Risk Factor Surveillance System)

- Nation's largest health survey (>400k respondents/year)
- Includes chronic disease diagnoses, behavior, demographic variables, and BMI
- Used for both *descriptive disparity analysis* and *predictive modeling*

From BRFSS, I constructed:

- A binary heart disease outcome
- Predictors including BMI, smoking status, hypertension, diabetes, physical activity
- Race variables consistent with CDC categories:
AIAN, Asian, Hispanic, NH_Black, NH_White, Other/Multi

Methodology

When data was obtained, it was then cleaned and processed. CDC weekly COVID mortality into aggregated annual rates. BRFSS variables were then filtered. Then created a binary target variable for heart disease, 0 for no, 1 for yes. In order to handle missing variables, BRFSS uses skip patterns.

My visuals are then created and tested through Python, and then organized through Tableau in two dashboards. First dashboard being “Public Health Disparities, including: COVID death rate trends(all adults), COVID death rate trends(50-64), BMI distribution, Average heart disease by race(2020). Second Dashboard is the “Predictive Model Fairness”, including: ROC curve(Logistic Regression), Race-stratified model fairness metrics(ACC, AUC, F1), Model comparison bar chart, Predictive fairness across race.

My baseline model consisted of logistic regression, stratified train/test split, standardization, and evaluation metrics: accuracy, AUC, and F1. Then, for the hyperparameter-tuned model, it completed a cross-validated grid search, penalty, C, solver tuning, and then selection based on AUC. The fairness evaluation across racial groups then occurs, for each racial group, it computes AUC, F1, and accuracy. This is then plotted and shown in Dashboard Two in Tableau to interpret disparities in the performance of the AI diagnosis model.

Dashboard 1: Public Health Disparities

Includes:

- COVID-19 death rate trends (all adults)
- COVID-19 death rate trends (50–64)
- BMI distribution (BRFSS 2022)
- Average heart disease prevalence by race (BRFSS 2020)

Dashboard 2: Predictive Model Fairness

Includes:

- ROC Curve (Logistic Regression)
- Race-stratified accuracy, AUC, and F1 metrics
- Model comparison: logistic regression, random forest, gradient boosting
- Fairness evaluation across race groups

Predictive Modeling Pipeline

1. 80/20 train-test split
2. Standardization of continuous predictors
3. Model training:
 - Logistic Regression
 - Random Forest
 - Gradient Boosting
4. Hyperparameter tuning:
 - GridSearchCV
 - 5-fold cross-validation
5. Evaluation Metrics:

- Accuracy
- AUC
- F1 Score (sensitive to class imbalance)

6. Race-Stratified Fairness Evaluation:

- Metrics computed separately for AIAN, Asian, Hispanic, NH Black, NH White

Results

Figures one and two show the COVID-19 mortality trend through loaded CDC data, then filtered by race and aggregated into annual summaries in Tableau. The “All adults” chart includes the entire U.S. population 18+. This heavily influences the older NH/White population. The “50-64 years” chart isolates a midlife band, where structural disadvantages dominate. These patterns are expected and crucial to interpreting disparities. Showing a clear racial connection in mortality. Groups such as NH/Black, AIAN, and Hispanic were the highest of the mortality trends. NH/White and Asian were showing lower mortality rates. The disparities persist in refined age group.

Dashboard one also consists of BMI Distribution shows a clear right shift for the entire population. This is a high-risk context for health outcome disparities, with each racial group having different possible outcomes. Also in this dashboard, there is Heart Disease Prevalence(BRFSS 2020) visual, showing the groups AIAN and NH/White as the highest and Asian as the lowest. Connecting the CDC mortality rate trend.

The age groups above and below make a difference to the delivered information. These metrics are all obtained from data from accounted for; impossible to tell what it could look like if there were more access to healthcare for lower-class minority racial groups. With the data obtained, we can see there is a slight disparity regardless of the population health disparity.

Epidemiological Findings (Dashboard 1)

COVID-19 Mortality Trends

- All adults (18+): Rates are dominated by older NH/White mortality due to population size
- Ages 50–64: NH/Black, Hispanic, and AIAN adults show the highest death rates

This demonstrates a key insight:

Aggregated statistics obscure disparities that appear clearly when age groups are separated.

Chronic Disease & BMI Disparities

- BMI distribution shows population-wide elevation, a risk factor for chronic disease
- Heart disease prevalence is highest among AIAN and NH/White adults
- Asian adults show consistently lower prevalence

These patterns reinforce structural disparities reflected in CDC mortality data.

Predictive Modeling Findings (Dashboard 2)

Overall Model Performance

- AUC $\approx 0.82\text{--}0.85$, showing good discrimination
- Logistic Regression selected as baseline due to interpretability

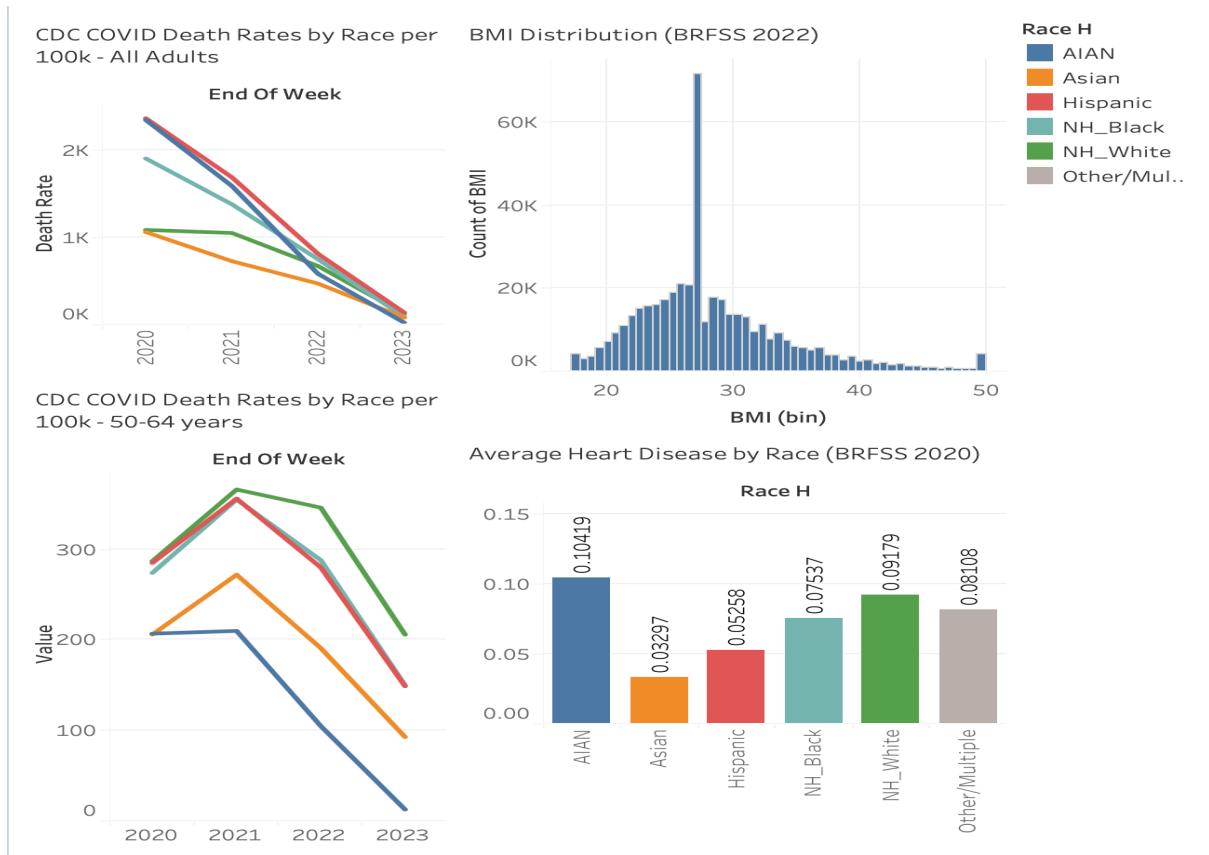
Race-Stratified Fairness

- Lower AUC for AIAN and NH_Black groups
- Higher AUC and accuracy for NH_White and Asian groups
- The F1 score is lower across minority groups due to class imbalance in BRFSS

This reveals a meaningful fairness concern:

Model accuracy and discrimination vary by race despite identical training workflows.

This indicates algorithmic bias rooted in underlying data inequalities, not the model architecture itself.



Predictive Model

The predictive model fairness dashboard represents the predictive model performance. This model was built using the BRFSS 2020 dataset, with complete heart disease predictor data. The model pipeline consisted of a train and test split of 80/20, then scaling continuous predictors, to model types tested, including logistic regression, random forest, and gradient boosting. A round of hyperparameter tuning via GridsearchCV, then a 5-fold cross-validation.

Finally, evaluation metrics for each model consisted of an accuracy, an AUC (area under the curve), and an F1 score. It contains three visuals: Model comparison, ROC Curve for logistic regression, and fairness by race. Representing logistic regression, random forest, and gradient boosting. By comparing ACC, AUC, and F1 for fairness evaluation. AUC was $\approx 0.82\text{--}0.85$ range showing good discrimination overall. Fairness AUC shows lower for NH/Black and AIAN. The F1 is more biased due to class imbalances. The model's performance is best for those in NH/White and Asian, showing an algorithmic disparity.



Discussion

Key findings in this project were that COVID and heart disease disparities are large and persistent and would show more accurately over more data. BRFSS health statistics and behaviors mirror these disparities. The logistic Regression models show uneven performance across racial groups, indicating algorithmic bias. This bias is shown in the structural inequality in real-world data, due to class imbalances, featured-race correlations, and missingness patterns in BRFSS. AI is not the bias; the bias is amplifying the disparities. Fairness-aware corrections that are needed would include reweighting, group-specific thresholds, and equalized odds constraints. Some limitations of this study. Although this is an informative study, there were some noticeable limitations, such as since the CDC data is aggregated and not individual-level, it's a broad number, and from some research, this is because not all people can be individually accounted for due to the volume of people. The BRFSS data is self-reported data, so there is no extra validation in the data from health professionals, which does affect the overall study validity. Missingness and selection bias exist in the data. Only one predictive model family was used, and it would be beneficial to use some others for more cross-validation.

Key findings:

- COVID-19 mortality disparities are large, persistent, and age-dependent
- BRFSS chronic disease and behavior patterns mirror these disparities

- The heart disease model shows uneven performance across racial groups, with minority groups receiving worse predictions
- Bias emerges from structural inequalities, feature-race correlations, and missingness patterns—not from the algorithm itself

Implications for AI in Healthcare

To prevent widening inequities, fairness-aware methods are essential:

- Reweighting underrepresented groups
- Group-specific thresholds
- Equalized-odds-based constraints
- Mandatory race-stratified validation before deployment

Limitations

- CDC data is aggregated, not individual-level
- BRFSS is self-reported, subject to recall bias
- Only one predictive model family (classical ML) was used
- Additional models (XGBoost, neural networks) could enrich the comparison

Conclusion

AI has the potential to expand healthcare access and improve clinical decision-making. However, without fairness checks, AI tools can unintentionally reinforce structural inequalities embedded in medical data.

This project shows that:

- Epidemiological disparities remain substantial across U.S. racial groups.
- Predictive models replicate these disparities unless monitored.

- Fairness auditing—AUC parity, equalized odds, race-stratified validation—is essential before deploying AI in clinical settings.

Ultimately, fair AI in healthcare is not automatic—it must be designed, tested, and regulated. Data scientists, clinicians, and policymakers must work collaboratively to ensure that AI advances do not widen racial inequities but instead contribute to addressing them.